

Relational Algebra

Basic Operations
Algebra of Bags

What is an “Algebra”

- ◆ Mathematical system consisting of:
 - ◆ *Operands* --- variables or values from which new values can be constructed.
 - ◆ *Operators* --- symbols denoting procedures that construct new values from given values.

What is Relational Algebra?

- ◆ An algebra whose operands are relations or variables that represent relations.
- ◆ Operators are designed to do the most common things that we need to do with relations in a database.
 - ◆ The result is an algebra that can be used as a *query language* for relations.

Core Relational Algebra

- ◆ Union, intersection, and difference.
 - ◆ Usual set operations, but *both operands must have the same relation schema.*
- ◆ Selection: picking certain rows.
- ◆ Projection: picking certain columns.
- ◆ Products and joins: compositions of relations.
- ◆ Renaming of relations and attributes.

Selection

◆ $R1 := \sigma_C(R2)$

- ◆ C is a condition (as in “if” statements) that refers to attributes of $R2$.
- ◆ $R1$ is all those tuples of $R2$ that satisfy C .

Example: Selection

Relation Sells:

| bar | beer | price |
|-------|--------|-------|
| Joe's | Bud | 2.50 |
| Joe's | Miller | 2.75 |
| Sue's | Bud | 2.50 |
| Sue's | Miller | 3.00 |

JoeMenu := $\sigma_{\text{bar}=\text{"Joe's"}}(\text{Sells})$:

| bar | beer | price |
|-------|--------|-------|
| Joe's | Bud | 2.50 |
| Joe's | Miller | 2.75 |

Projection

◆ $R1 := \pi_L(R2)$

- ◆ L is a list of attributes from the schema of $R2$.
- ◆ $R1$ is constructed by looking at each tuple of $R2$, extracting the attributes on list L , in the order specified, and creating from those components a tuple for $R1$.
- ◆ Eliminate duplicate tuples, if any.

Example: Projection

Relation Sells:

| bar | beer | price |
|-------|--------|-------|
| Joe's | Bud | 2.50 |
| Joe's | Miller | 2.75 |
| Sue's | Bud | 2.50 |
| Sue's | Miller | 3.00 |

Prices := $\pi_{\text{beer, price}}(\text{Sells})$:

| beer | price |
|--------|-------|
| Bud | 2.50 |
| Miller | 2.75 |
| Miller | 3.00 |

Extended Projection

- ◆ Using the same π_L operator, we allow the list L to contain arbitrary expressions involving attributes:
 1. Arithmetic on attributes, e.g., $A + B \rightarrow C$.
 2. Duplicate occurrences of the same attribute.

Example: Extended Projection

$R =$ (

| A | B |
|---|---|
| 1 | 2 |
| 3 | 4 |

)

$\pi_{A+B \rightarrow C, A, A}(R) =$

| C | A1 | A2 |
|---|----|----|
| 3 | 1 | 1 |
| 7 | 3 | 3 |

Product

◆ $R3 := R1 \times R2$

- ◆ Pair each tuple $t1$ of $R1$ with each tuple $t2$ of $R2$.
- ◆ Concatenation $t1t2$ is a tuple of $R3$.
- ◆ Schema of $R3$ is the attributes of $R1$ and then $R2$, in order.
- ◆ But beware attribute A of the same name in $R1$ and $R2$: use $R1.A$ and $R2.A$.

Example: $R3 := R1 \times R2$

R1(

| A, | B) |
|----|----|
| 1 | 2 |
| 3 | 4 |

R2(

| B, | C) |
|----|----|
| 5 | 6 |
| 7 | 8 |
| 9 | 10 |

R3(

| A, | R1.B, | R2.B, | C) |
|----|-------|-------|----|
| 1 | 2 | 5 | 6 |
| 1 | 2 | 7 | 8 |
| 1 | 2 | 9 | 10 |
| 3 | 4 | 5 | 6 |
| 3 | 4 | 7 | 8 |
| 3 | 4 | 9 | 10 |

Theta-Join

- ◆ $R3 := R1 \bowtie_C R2$
 - ◆ Take the product $R1 \times R2$.
 - ◆ Then apply σ_C to the result.
- ◆ As for σ , C can be any boolean-valued condition.
 - ◆ Historic versions of this operator allowed only $A \theta B$, where θ is $=$, $<$, etc.; hence the name “theta-join.”

Example: Theta Join

Sells(

| bar, | beer, | price |
|-------|--------|-------|
| Joe's | Bud | 2.50 |
| Joe's | Miller | 2.75 |
| Sue's | Bud | 2.50 |
| Sue's | Coors | 3.00 |

)

Bars(

| name, | addr |
|-------|-----------|
| Joe's | Maple St. |
| Sue's | River Rd. |

)

BarInfo := Sells $\bowtie_{\text{Sells.bar} = \text{Bars.name}}$ Bars

BarInfo(

| bar, | beer, | price, | name, | addr |
|-------|--------|--------|-------|-----------|
| Joe's | Bud | 2.50 | Joe's | Maple St. |
| Joe's | Miller | 2.75 | Joe's | Maple St. |
| Sue's | Bud | 2.50 | Sue's | River Rd. |
| Sue's | Coors | 3.00 | Sue's | River Rd. |

)

Natural Join

- ◆ A useful join variant (*natural* join) connects two relations by:
 - ◆ Equating attributes of the same name, and
 - ◆ Projecting out one copy of each pair of equated attributes.
- ◆ Denoted $R3 := R1 \bowtie R2$.

Example: Natural Join

Sells(bar, beer, price)

| bar | beer | price |
|-------|--------|-------|
| Joe's | Bud | 2.50 |
| Joe's | Miller | 2.75 |
| Sue's | Bud | 2.50 |
| Sue's | Coors | 3.00 |

Bars(bar, addr)

| bar | addr |
|-------|-----------|
| Joe's | Maple St. |
| Sue's | River Rd. |

BarInfo := Sells ⋈ Bars

Note: Bars.name has become Bars.bar to make the natural join “work.”

BarInfo(bar, beer, price, addr)

| bar | beer | price | addr |
|-------|--------|-------|-----------|
| Joe's | Bud | 2.50 | Maple St. |
| Joe's | Miller | 2.75 | Maple St. |
| Sue's | Bud | 2.50 | River Rd. |
| Sue's | Coors | 3.00 | River Rd. |

Renaming

- ◆ The ρ operator gives a new schema to a relation.
- ◆ $R1 := \rho_{R1(A1, \dots, An)}(R2)$ makes R1 be a relation with attributes $A1, \dots, An$ and the same tuples as R2.
- ◆ Simplified notation: $R1(A1, \dots, An) := R2$.

Example: Renaming

Bars(

| name, | addr |
|-------|-----------|
| Joe's | Maple St. |
| Sue's | River Rd. |

)

$R(\text{bar}, \text{addr}) := \text{Bars}$

R(

| bar, | addr |
|-------|-----------|
| Joe's | Maple St. |
| Sue's | River Rd. |

)

Building Complex Expressions

- ◆ Combine operators with parentheses and precedence rules.
- ◆ Three notations, just as in arithmetic:
 1. Sequences of assignment statements.
 2. Expressions with several operators.
 3. Expression trees.

Sequences of Assignments

- ◆ Create temporary relation names.
- ◆ Renaming can be implied by giving relations a list of attributes.
- ◆ **Example:** $R3 := R1 \bowtie_C R2$ can be written:
 $R4 := R1 \times R2$
 $R3 := \sigma_C(R4)$

Expressions in a Single Assignment

- ◆ **Example:** the theta-join $R3 := R1 \bowtie_c R2$
can be written: $R3 := \sigma_c(R1 \times R2)$
- ◆ Precedence of relational operators:
 1. $[\sigma, \pi, \rho]$ (highest).
 2. $[\times, \bowtie]$.
 3. \cap .
 4. $[\cup, -]$

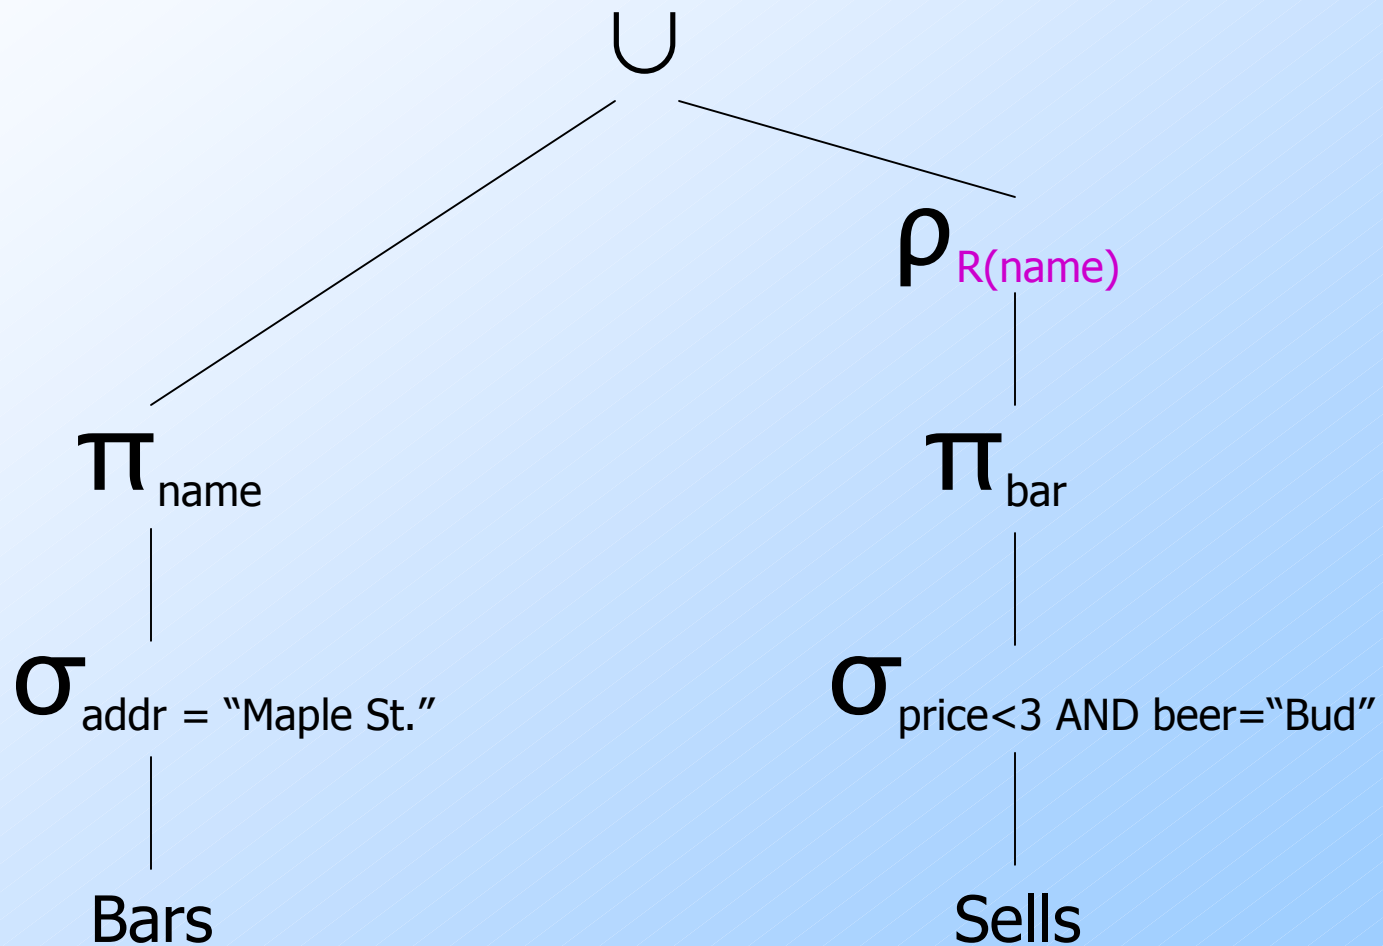
Expression Trees

- ◆ Leaves are operands --- either variables standing for relations or particular, constant relations.
- ◆ Interior nodes are operators, applied to their child or children.

Example: Tree for a Query

- ◆ Using the relations **Bars(name, addr)** and **Sells(bar, beer, price)**, find the names of all the bars that are either on Maple St. or sell Bud for less than \$3.

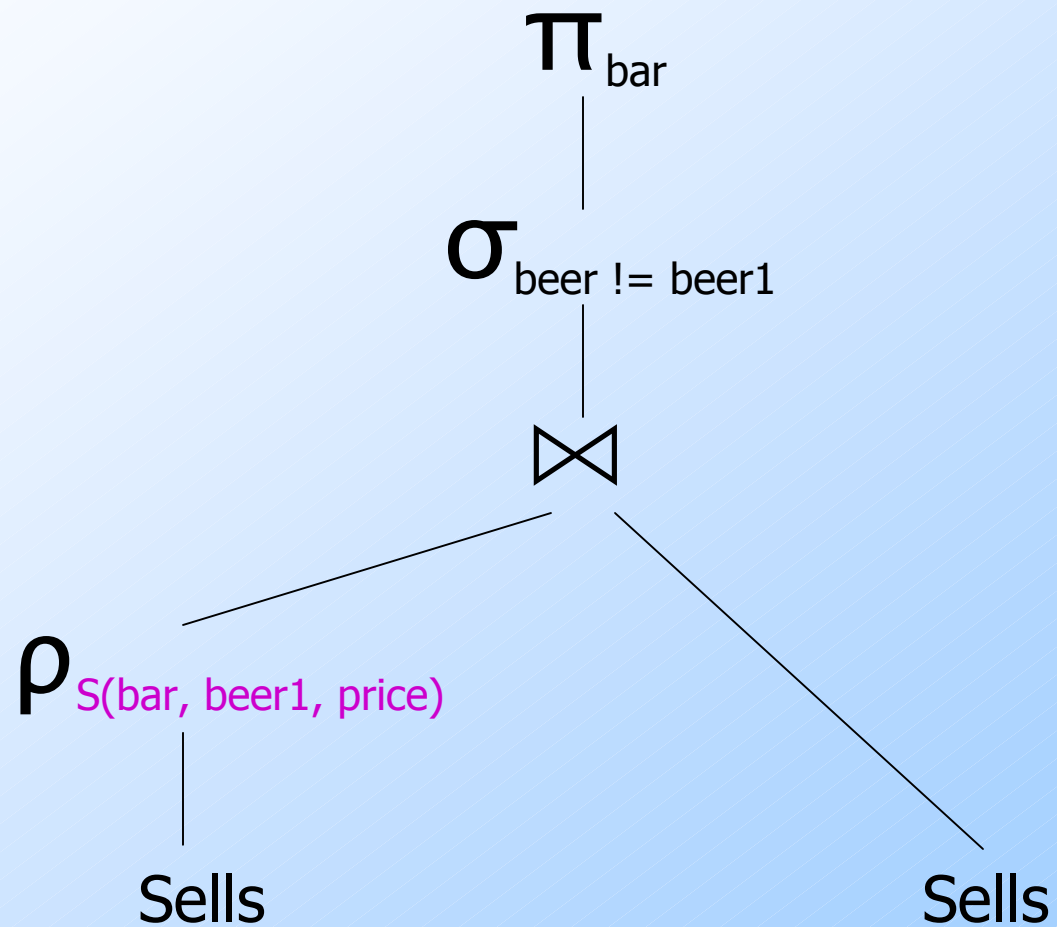
As a Tree:



Example: Self-Join

- ◆ Using $\text{Sells}(\text{bar}, \text{beer}, \text{price})$, find the bars that sell two different beers at the same price.
- ◆ **Strategy**: by renaming, define a copy of Sells , called $\text{S}(\text{bar}, \text{beer1}, \text{price})$. The natural join of Sells and S consists of quadruples $(\text{bar}, \text{beer}, \text{beer1}, \text{price})$ such that the bar sells both beers at this price.

The Tree



Schemas for Results

- ◆ **Union, intersection, and difference:** the schemas of the two operands must be the same, so use that schema for the result.
- ◆ **Selection:** schema of the result is the same as the schema of the operand.
- ◆ **Projection:** list of attributes tells us the schema.

Schemas for Results --- (2)

- ◆ **Product**: schema is the attributes of both relations.
 - ◆ Use $R.A$, etc., to distinguish two attributes named A .
- ◆ **Theta-join**: same as product. since selection does not change schema
- ◆ **Natural join**: union of the attributes of the two relations.
- ◆ **Renaming**: the operator tells the schema.

Relational Algebra on Bags

- ◆ A *bag* (or *multiset*) is like a set, but an element may appear more than once.

tuple of the relation

- ◆ Example: $\{1,2,1,3\}$ is a bag.

- ◆ Example: $\{1,2,3\}$ is also a bag that happens to be a set.

Why Bags?

- ◆ SQL, the most important query language for relational databases, is actually a bag language.
- ◆ Some operations, like projection, are more efficient on bags than sets.
since no need to eliminate duplicates

Operations on Bags

selection keeps duplicates

- ◆ **Selection** applies to each tuple, so its effect on bags is like its effect on sets.

projection keeps duplicates

- ◆ **Projection** also applies to each tuple, but as a bag operator, we do not eliminate duplicates.

- ◆ **Products** and **joins** are done on each pair of tuples, so duplicates in bags have no effect on how we operate.

product and join keep duplicates

Example: Bag Selection

R(

| A, | B |
|----|---|
| 1 | 2 |
| 5 | 6 |
| 1 | 2 |

)

$\sigma_{A+B < 5} (R) =$

| A | B |
|---|---|
| 1 | 2 |
| 1 | 2 |

Example: Bag Projection

R(

| A, | B |
|----|---|
| 1 | 2 |
| 5 | 6 |
| 1 | 2 |

)

$\pi_A(R) =$

| A |
|---|
| 1 |
| 5 |
| 1 |

Example: Bag Product

R(

| A, | B |
|----|---|
| 1 | 2 |
| 5 | 6 |
| 1 | 2 |

)

S(

| B, | C |
|----|---|
| 3 | 4 |
| 7 | 8 |

)

R X S =

| A | R.B | S.B | C |
|---|-----|-----|---|
| 1 | 2 | 3 | 4 |
| 1 | 2 | 7 | 8 |
| 5 | 6 | 3 | 4 |
| 5 | 6 | 7 | 8 |
| 1 | 2 | 3 | 4 |
| 1 | 2 | 7 | 8 |

Example: Bag Theta-Join

R(

| A, | B |
|----|---|
| 1 | 2 |
| 5 | 6 |
| 1 | 2 |

)

S(

| B, | C |
|----|---|
| 3 | 4 |
| 7 | 8 |

)

R $\bowtie_{R.B < S.B}$ S =

| A | R.B | S.B | C |
|---|-----|-----|---|
| 1 | 2 | 3 | 4 |
| 1 | 2 | 7 | 8 |
| 5 | 6 | 7 | 8 |
| 1 | 2 | 3 | 4 |
| 1 | 2 | 7 | 8 |

Bag Union

- ◆ An element appears in the union of two bags the sum of the number of times it appears in each bag.
- ◆ **Example:** $\{1,2,1\} \cup \{1,1,2,3,1\} = \{1,1,1,1,1,2,2,3\}$

Bag Intersection

tuple of the relation



- ◆ An element appears in the intersection of two bags the minimum of the number of times it appears in either.
- ◆ **Example:** $\{1,2,1,1\} \cap \{1,2,1,3\} = \{1,1,2\}$.

Bag Difference

tuple of the relation

- ◆ An element appears in the difference $A - B$ of bags as many times as it appears in A , minus the number of times it appears in B .
 - ◆ But never less than 0 times.
- ◆ **Example:** $\{1,2,1,1\} - \{1,2,3\} = \{1,1\}$.

Beware: Bag Laws \neq Set Laws

- ◆ Some, but *not all* algebraic laws that hold for sets also hold for bags.
- ◆ **Example:** the commutative law for union ($R \cup S = S \cup R$) *does* hold for bags.
 - ◆ Since addition is commutative, adding the number of times x appears in R and S doesn't depend on the order of R and S .

Example: A Law That Fails

- ◆ Set union is *idempotent*, meaning that $S \cup S = S$.
- ◆ However, for bags, if x appears n times in S , then it appears $2n$ times in $S \cup S$.
- ◆ Thus $S \cup S \neq S$ in general.
 - ◆ e.g., $\{1\} \cup \{1\} = \{1,1\} \neq \{1\}$.