

Transactions: Basic Definition

A transaction (“TXN”) is a sequence of one or more *operations* (reads or writes) which reflects *a single real-world transition*.

In the real world, a TXN either happened completely or not at all

```
START TRANSACTION;  
    UPDATE Bank SET amount = amount - 100  
    WHERE name = 'Bob';  
    UPDATE Bank SET amount = amount + 100  
    WHERE name = 'Joe';  
COMMIT; (OR ROLLBACK;)
```

Motivation for Transactions

Grouping user actions (reads & writes) into *transactions* helps with two goals:

Recovery & Durability: Keeping the DBMS data consistent and durable in the face of crashes, aborts, system shutdowns, etc.

Concurrency: Achieving better performance by parallelizing TXNs *without* creating anomalies

Motivation

1. Recovery & Durability of user data is essential for reliable DBMS usage

- The DBMS may experience crashes (e.g. power outages, etc.)
- Individual TXNs may be aborted (e.g. by the user)

Idea: Make sure that TXNs are either **durably stored in full, or not at all**; keep log to be able to “roll-back” TXNs

Protection against crashes / aborts

Client 1:

```
INSERT INTO SmallProduct(name, price)
SELECT pname, price
FROM Product
WHERE price <= 0.99;
```

Crash / abort!

```
DELETE Product
WHERE price <= 0.99;
```

What goes wrong?

Protection against crashes / aborts

Client 1:

```
START TRANSACTION;  
  INSERT INTO SmallProduct(name, price)  
  SELECT pname, price  
  FROM Product  
  WHERE price <= 0.99;  
  
  DELETE Product  
  WHERE price <= 0.99;  
COMMIT;
```

Atomicity & Durability

- **Atomicity:**

- TXNs should either happen completely or not at all
- If abort / crash during TXN, *no* effects should be seen

TXN 1



Crash / abort

No changes
persisted

- **Durability:**

- If DBMS stops running, changes due to completed TXNs should all persist
- *Just store on stable disk*

TXN 2



All changes
persisted

ACID: Atomicity

- TXN's activities are atomic: **all or nothing**
 - Intuitively: in the real world, a transaction is something that would either occur *completely* or *not at all*
- Two possible outcomes for a TXN
 - It *commits*: all the changes are made
 - It *aborts*: no changes are made

ACID: Durability

- The effect of a TXN must continue to exist (“***persist***”) after the TXN
 - And after the whole program has terminated
 - And even if there are power failures, crashes, etc.
 - And etc...
- Means: Write data to **disk**

Motivation

2. Concurrent execution of user programs is essential for good DBMS performance.

- Individual TXNs might be *slow*- don't want to block other users during
- Disk access may be *slow*- let some TXNs use CPUs while others accessing disk
- Users should still be able to execute TXNs as if in **isolation** and such that **consistency** is maintained

Idea: Have the DBMS handle running several user TXNs concurrently, in order to keep CPUs humming...

Example- consider two TXNs:

```
T1: START TRANSACTION;  
    UPDATE Accounts  
    SET Amt = Amt + 100  
    WHERE Name = 'A';  
  
    UPDATE Accounts  
    SET Amt = Amt - 100  
    WHERE Name = 'B';  
COMMIT;
```

T1 transfers \$100 from B's account to A's account

```
T2: START TRANSACTION;  
    UPDATE Accounts  
    SET Amt = Amt * 1.06;  
COMMIT;
```


T2 credits both accounts with a 6% interest payment

Scheduling examples

*Starting
Balance*


A	B
\$50	\$200

Serial schedule T_1, T_2 :



A	B
\$159	\$106

Interleaved schedule A:



A	B
\$159	\$106


Same
result!

Scheduling examples

Starting Balance


A	B
\$50	\$200

Serial schedule T_1, T_2 :



A	B
\$159	\$106

Interleaved schedule B:



A	B
\$159	\$112


Different result than serial T_1, T_2 !

Scheduling examples

Starting
Balance


A	B
\$50	\$200

Serial schedule T_2, T_1 :



A	B
\$153	\$112

Interleaved schedule B:



A	B
\$159	\$112

Different
result than
serial
 T_2, T_1
ALSO!

Concurrency: Isolation & Consistency

- The DBMS must handle concurrency such that...

Isolation is maintained: Users must be able to execute each TXN **as if they were the only user**

- DBMS handles the details of *interleaving* various TXNs

Consistency is maintained: TXNs must leave the DB in a **consistent state**

- DBMS handles the details of enforcing integrity constraints

ACID: Consistency

- The tables must always satisfy user-specified ***integrity constraints***
 - *Examples:*
 - Account number is unique
 - Stock amount can't be negative
- How consistency is achieved:
 - Programmer makes sure a txn takes a consistent state to a consistent state
 - *System* makes sure that the txn is **atomic**

ACID: Isolation

- A transaction executes concurrently with other transactions
- **Isolation**: the effect is as if each transaction executes in *isolation* of the others.
 - E.g. Should not be able to observe changes from other transactions during the run

Isolation Levels

- Fully serializable isolation is more expensive than “no isolation”
 - We can't do as many things concurrently
- For performance, we generally want to specify the most relaxed *isolation level* that's acceptable

Choosing the Isolation Level

- Before a transaction, we can say:

SET TRANSACTION ISOLATION LEVEL X

where X =

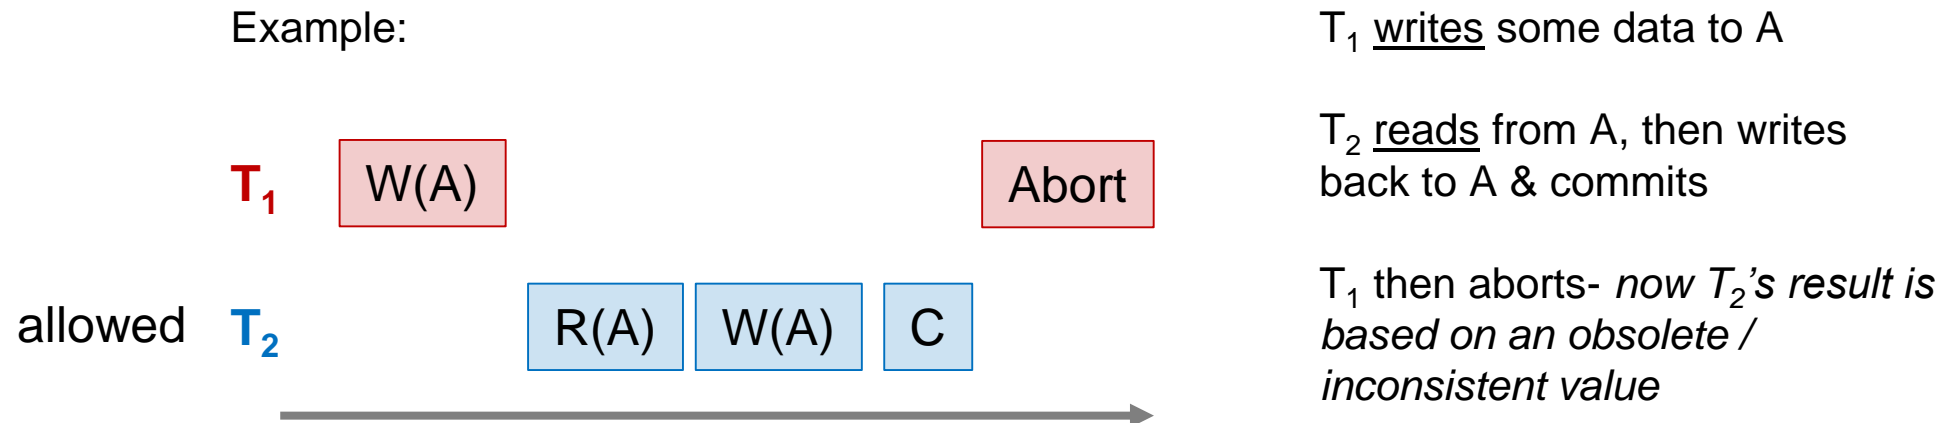
1. SERIALIZABLE (Default)
 2. REPEATABLE READ
 3. READ COMMITTED
 4. READ UNCOMMITTED
- This affects what data the next transaction may see

4 Read Uncommitted (T_2)

- A transaction running under READ UNCOMMITTED can see data in the database, even if it was written by a transaction that has not committed (and may never).

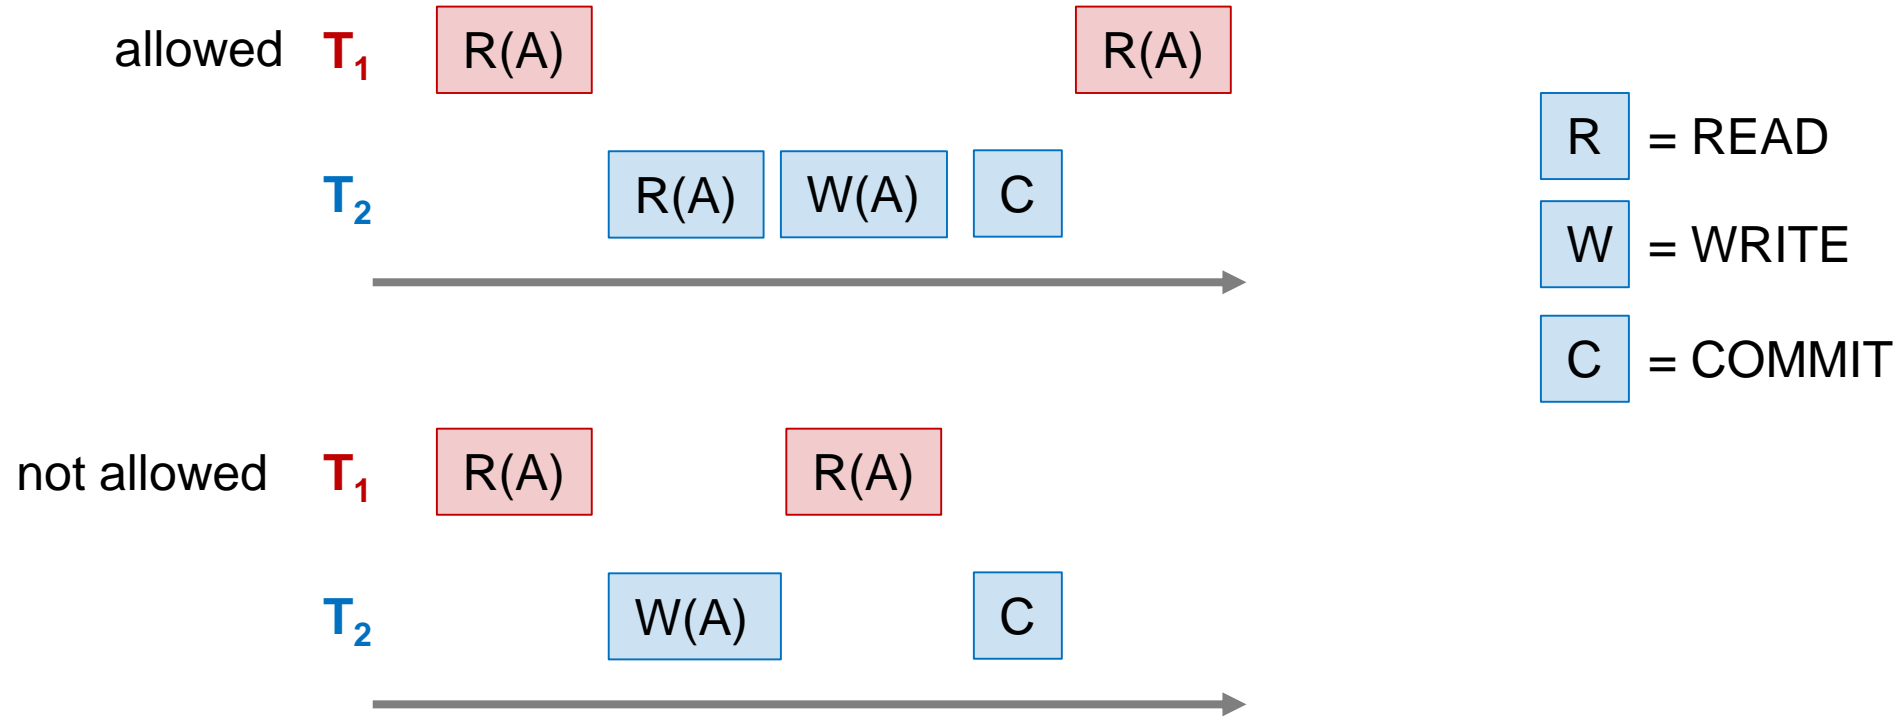
“Dirty read” / Reading uncommitted data:

Example:



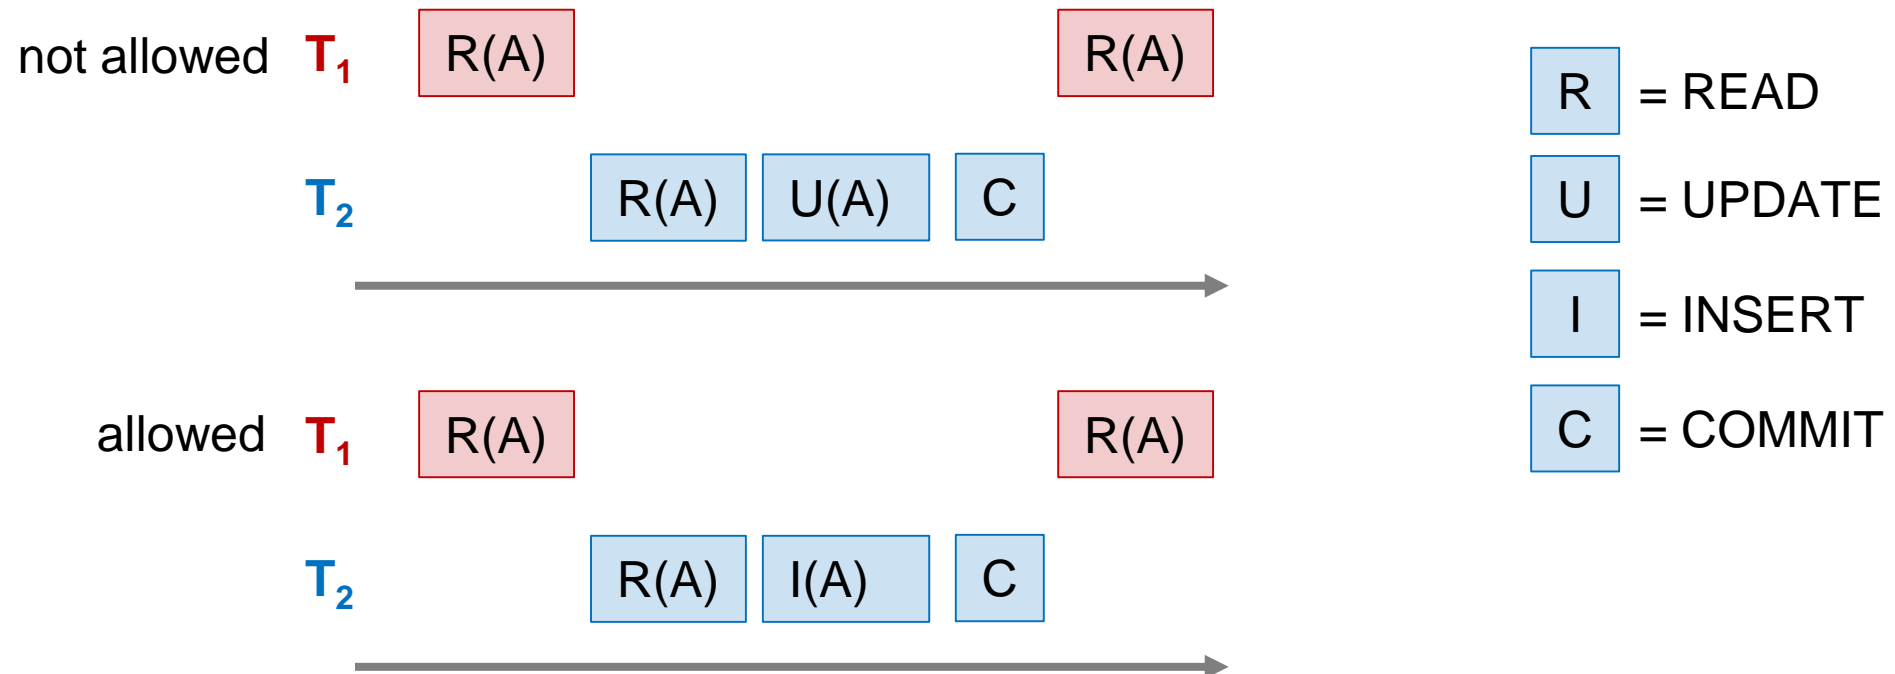
3 Read Committed (T_1)

- can see only committed data, but not necessarily the same data each time

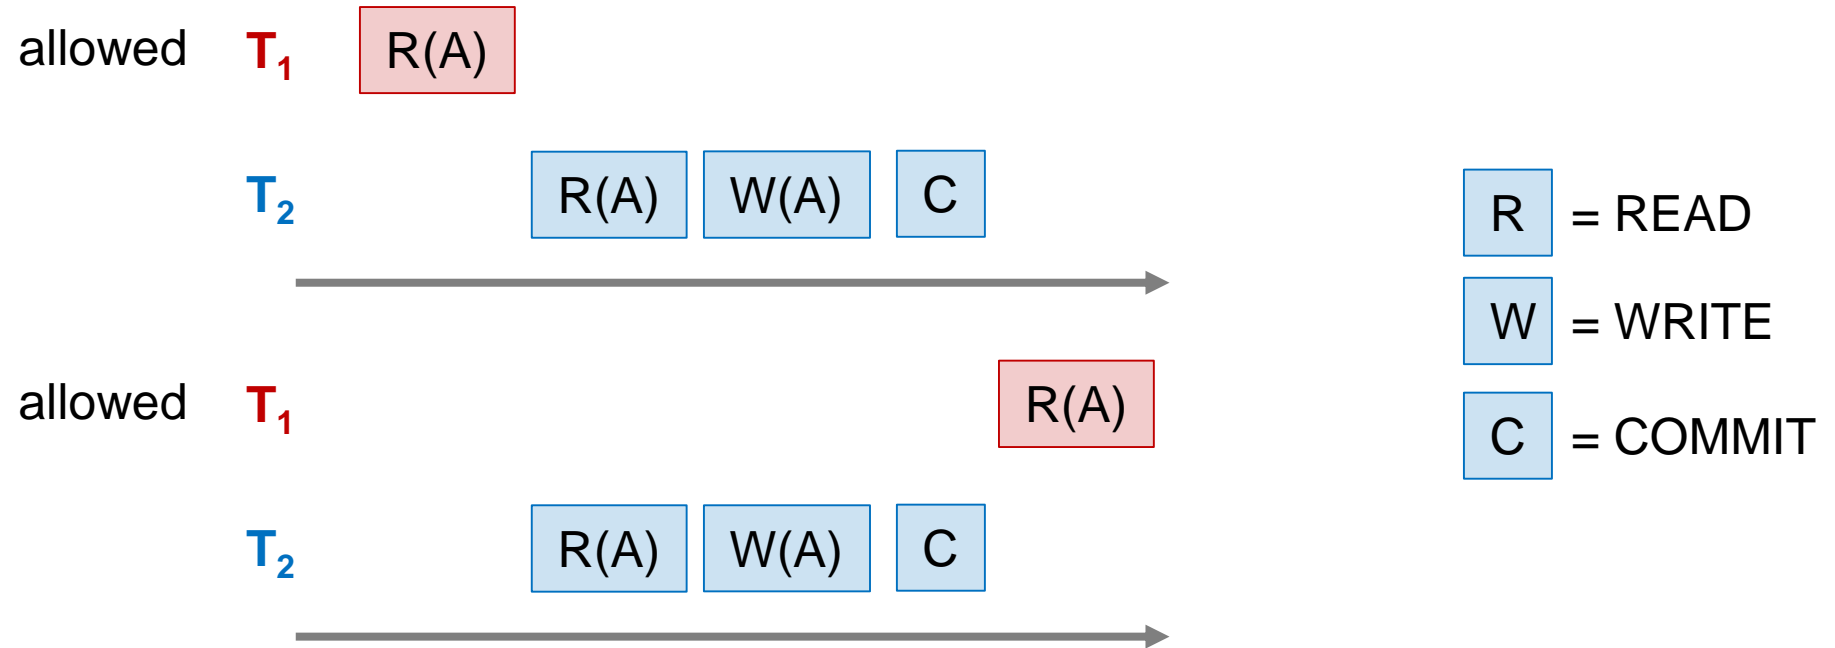


2 Repeatable Read (T_1)

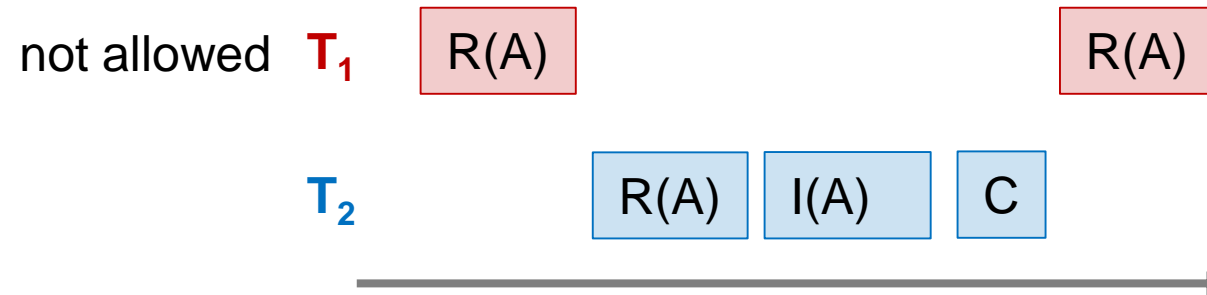
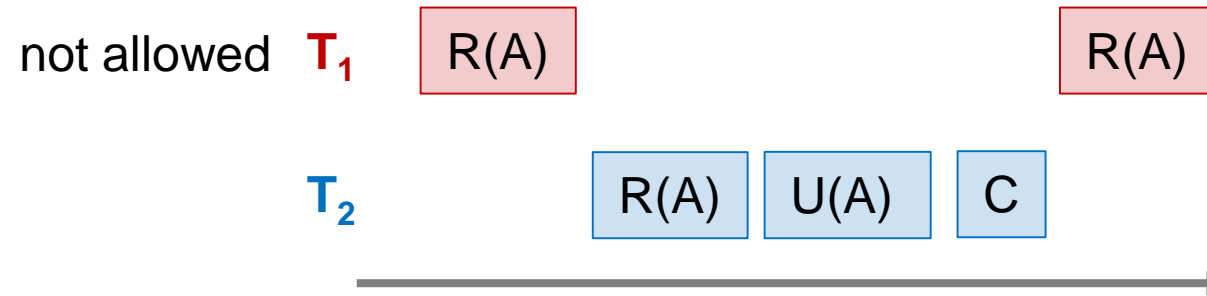
- Requirement is like read committed, plus: if data is read again, then everything seen the first time will be seen the second time.
 - But the second and subsequent reads may see *more* tuples as well.



1 Serializable (T_1)



1 Serializable (T_1)



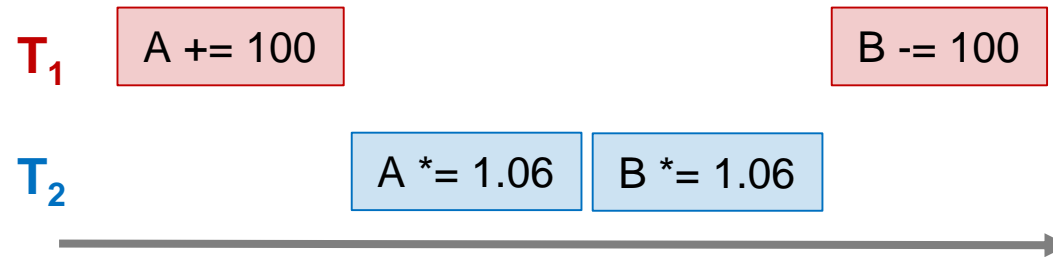
R = READ
U = UPDATE
I = INSERT
C = COMMIT

Scheduling Definitions

- A **serial schedule** is one that does not interleave the actions of different transactions
- A and B are **equivalent schedules** if, *for any database state*, the effect on DB of executing A **is identical to** the effect of executing B
- A **serializable schedule** is a schedule that is equivalent to **some** serial execution of the transactions.

Scheduling examples

Interleaved schedule B:



This schedule is different than ***any serial order!*** We say that it is **not serializable**

A	B
\$159	\$112


A	B
\$159	\$106

T_1, T_2


A	B
\$153	\$112

T_2, T_1

Serial Schedules




S1




S2


Interleaved Schedules




S3



S4



S5



S6

Serial Schedules	S1, S2
Serializable Schedules	S3, S4
Equivalent Schedules	<S1, S3> <S2, S4>
Non-serializable (Bad) Schedules	S5, S6

Conflict Types

Two actions of two different TXNs **conflict** if they involve the same variable, and at least one of them is a write

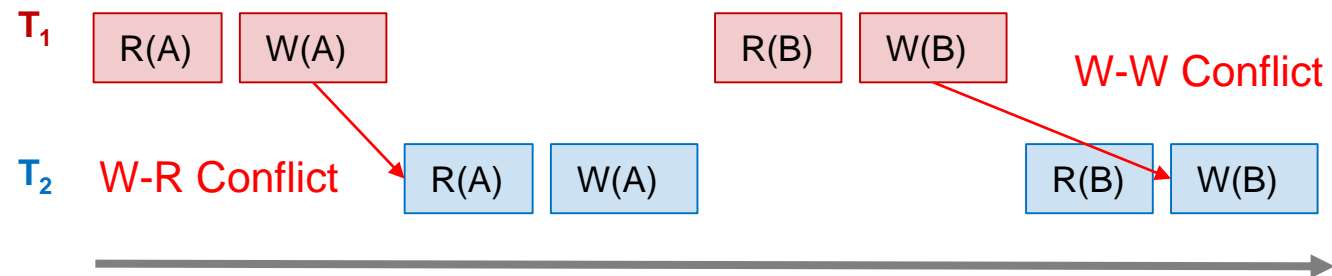
Thus, there are three types of conflicts:

- Read-Write conflicts (RW)
- Write-Read conflicts (WR)
- Write-Write conflicts (WW)

If we interchange the order of two conflicting actions, then the behavior of at least one TXN can change

Conflicts

Two actions of two different TXNs **conflict** if they involve the same variable, and at least one of them is a write



If we interchange the order of two conflicting actions, then the behavior of at least one TXN can change

Conflict Serializability

Two schedules are ***conflict equivalent*** if:

- They involve *the same actions of the same TXNs*
- Every *pair of conflicting actions* of two TXNs are *ordered in the same way*


Schedule S is ***conflict serializable*** if S is *conflict equivalent* to ***some*** serial schedule

Conflict serializable \Rightarrow serializable

So if we have conflict serializable, we have consistency & isolation!


Example “Good” vs. “bad” schedules

Serial Schedule:




Note that in the “bad” schedule, the **order of conflicting actions is different than the above (or any) serial schedule!**

Interleaved Schedules:




The Conflict Graph

- Let's now consider looking at conflicts **at the TXN level**
- Consider a graph where the **nodes are TXNs**, and there is an edge from $T_i \rightarrow T_j$ **if any actions in T_i precede and conflict with any actions in T_j**




What can we say about “good” vs. “bad” conflict graphs?

Serial Schedule:




A bit complicated...

Interleaved Schedules:




What can we say about “good” vs.
“bad” conflict graphs?

Serial Schedule:



Simple!

Interleaved Schedules:



Theorem: Schedule is **conflict serializable** if
and only if its conflict graph is **acyclic**